HiFormer–TencentGR: attention redistribution

Same checkpoint · seed 20262028 · target attention · head 0

Low-TV case (P10) · TV = 0.1182 · 90/100 valid keys · target item 8846555

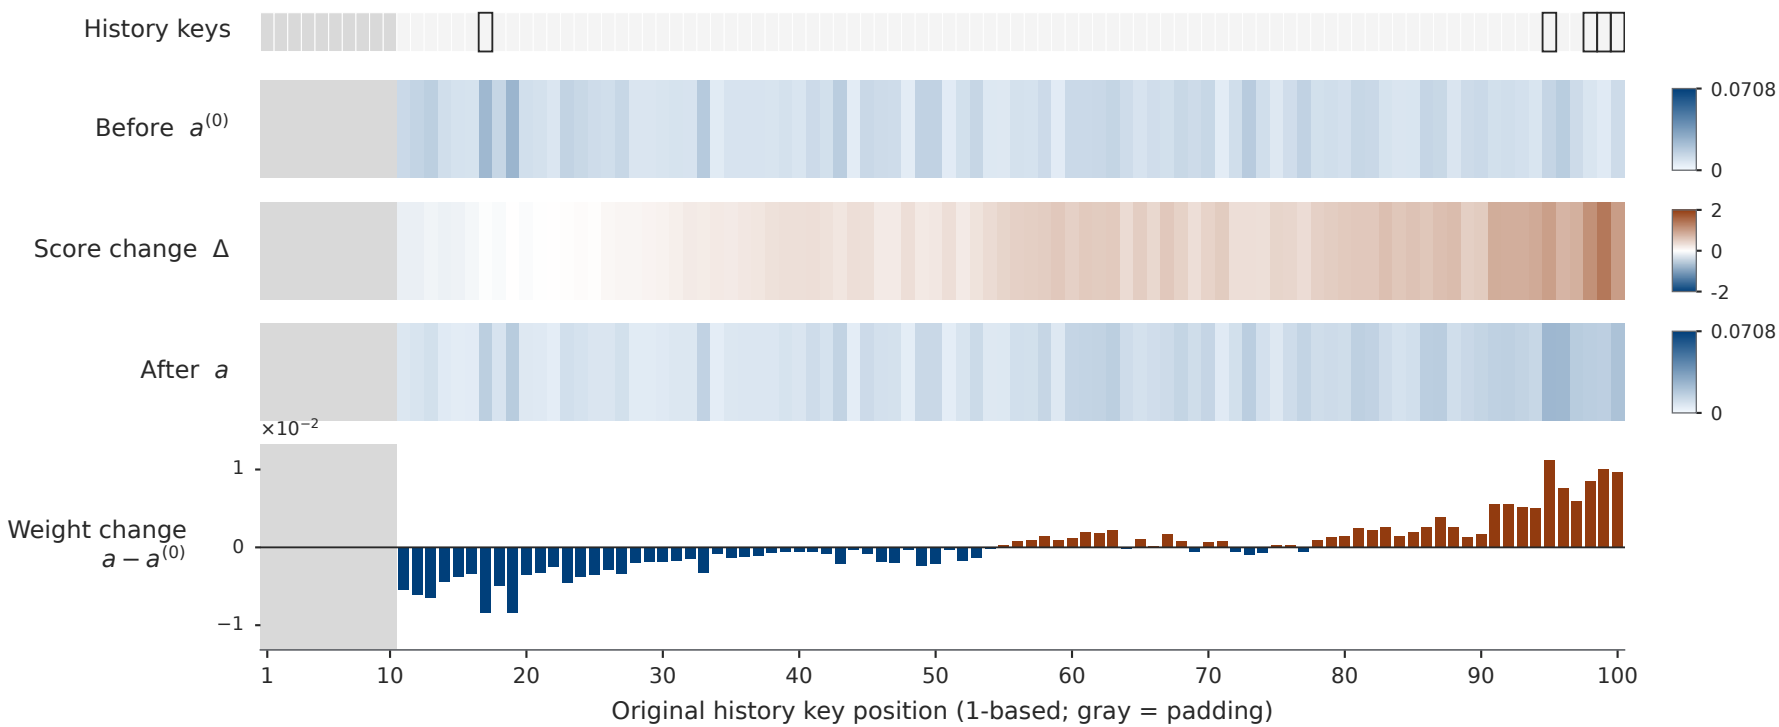

Position annotations: five largest absolute weight changes within this selected case

Position	History item ID	History action	Δ	$a^{(0)}$	a	$a - a^{(0)}$
17	9392554	conversion (2)	-0.0237	0.02489	0.01645	-0.00844
95	7533737	exposure (3)	+0.9766	0.01401	0.02517	+0.01116
98	7546248	exposure (3)	+1.1094	0.00809	0.01660	+0.00851
99	11478991	exposure (3)	+1.3672	0.00608	0.01616	+0.01007
100	5153566	exposure (3)	+0.9922	0.01166	0.02128	+0.00962

Sample 429 · part 6 · row 12230 · case chosen by TV quantile, without prediction/label selection.

Same-checkpoint local comparison, not an independently trained baseline. FP32 reconstruction; full allowed-key softmax; no re-normalization.

Before/after share a color scale; score colors are zero-centered. Item IDs are anonymous; key position does not imply elapsed time.

History action vocabulary: click=1, conversion=2, exposure=3. Target action 0 is not interpreted as a behavior.