HiFormer–TencentGR: attention redistribution

Same checkpoint · seed 20262028 · target attention · head 0

High-TV case (P90) · TV = 0.2101 · 70/100 valid keys · target item 6913384

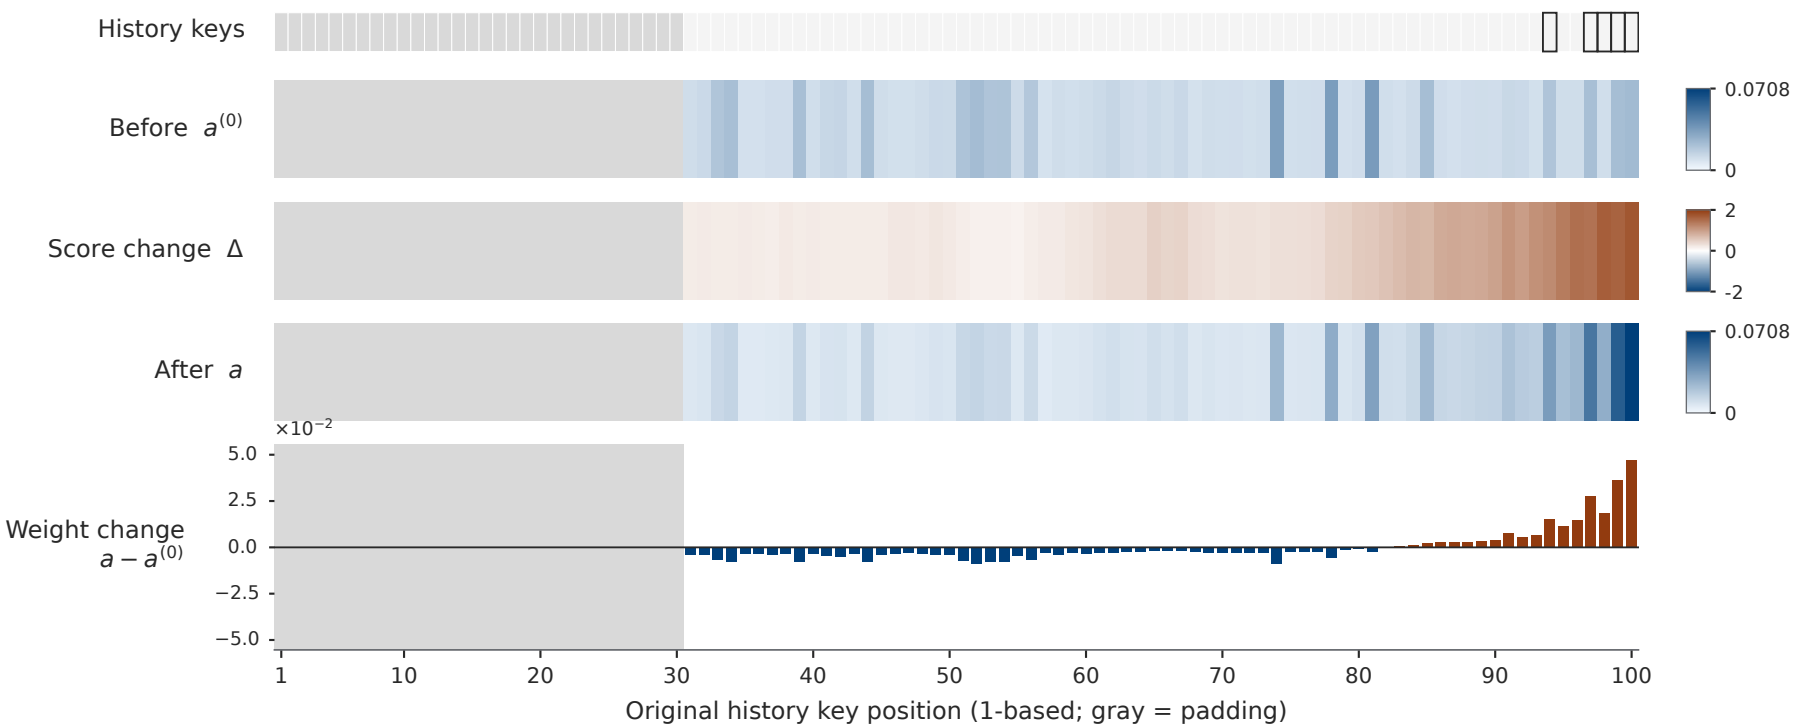

Position annotations: five largest absolute weight changes within this selected case

Position	History item ID	History action	Δ	$a^{(0)}$	a	$a - a^{(0)}$
94	8113026	exposure (3)	+1.1797	0.02045	0.03565	+0.01520
97	703267	exposure (3)	+1.4297	0.02224	0.04977	+0.02754
98	11059894	exposure (3)	+1.6250	0.01059	0.02881	+0.01822
99	11784798	exposure (3)	+1.5781	0.02280	0.05920	+0.03640
100	3135887	exposure (3)	+1.7109	0.02387	0.07077	+0.04690

Sample 150 · part 2 · row 5025 · case chosen by TV quantile, without prediction/label selection.

Same-checkpoint local comparison, not an independently trained baseline. FP32 reconstruction; full allowed-key softmax; no re-normalization.

Before/after share a color scale; score colors are zero-centered. Item IDs are anonymous; key position does not imply elapsed time.

History action vocabulary: click=1, conversion=2, exposure=3. Target action 0 is not interpreted as a behavior.