HiFormer–TencentGR: attention redistribution

Same checkpoint · seed 20262028 · target attention · head 0

Median-TV case (P50) · TV = 0.1630 · 92/100 valid keys · target item 5126221

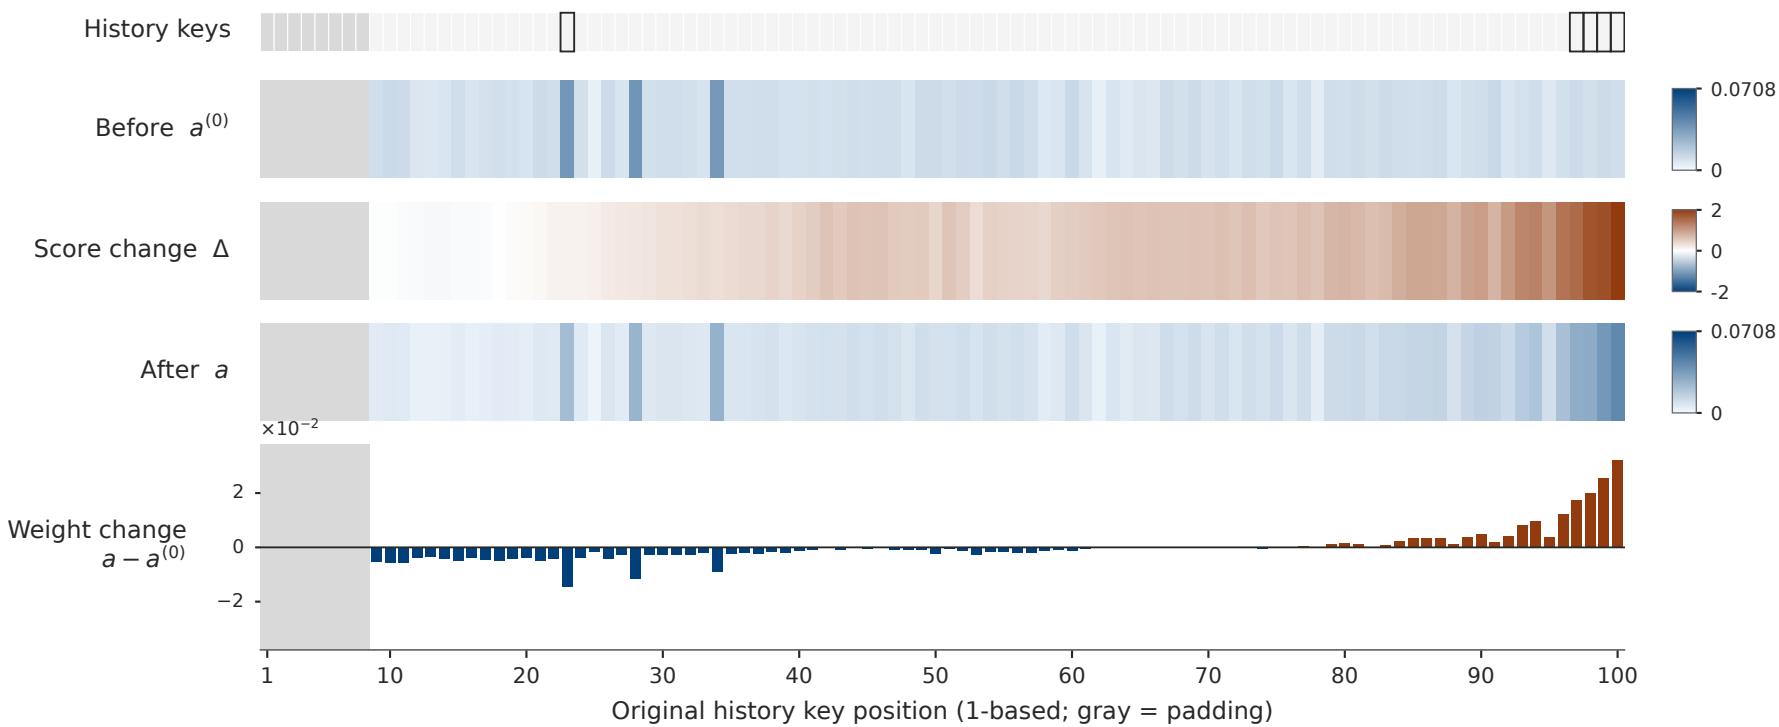

Position annotations: five largest absolute weight changes within this selected case

Position	History item ID	History action	Δ	$a^{(0)}$	a	$a - a^{(0)}$
23	4295713	click (1)	+0.1299	0.03759	0.02301	-0.01457
97	14390345	exposure (3)	+1.5078	0.01216	0.02954	+0.01738
98	15854583	exposure (3)	+1.7188	0.01002	0.03006	+0.02004
99	5703651	exposure (3)	+1.7812	0.01164	0.03716	+0.02552
100	99146	exposure (3)	+2.0000	0.01077	0.04279	+0.03202

Sample 128 · part 2 · row 413 · case chosen by TV quantile, without prediction/label selection.

Same-checkpoint local comparison, not an independently trained baseline. FP32 reconstruction; full allowed-key softmax; no re-normalization.

Before/after share a color scale; score colors are zero-centered. Item IDs are anonymous; key position does not imply elapsed time.

History action vocabulary: click=1, conversion=2, exposure=3. Target action 0 is not interpreted as a behavior.